

Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur
CS69202: Software Engineering Lab, Spring 2023
Assignment – 3

Submission Date – 28th February 2023

Question (60 + 30 + 10 = 100 marks):

Write a C++ program for implementing some functionalities of a Relational database using basic operators from Relational Algebra.

A relational database is a database where all data are organised in terms of *relations* or *tables* of *records* or *tuples* or *rows*. Relations can be thought of as sets of records. Each record is an array of *attribute* values (also called *columns*). Finally, the database is a collection of such relations. The attributes in a relation are also described in a *schema* using attribute names, which are used as keys for various operations on the data in a database. For more details, see:

https://en.wikipedia.org/wiki/Relational_database

In your program, implement the following classes:

```
class Attr { //Base class for attributes
    // Add operators for different attribute type in derived classes
    public: virtual bool operator==(const Attr & right) = 0;
    ...
    ...
    ...
};

class Record { //storing data for each record
    vector<Attr*> attrptr;
    //methods
    ...
    ...
};

class Relation { // storing a relation
    int nattr, nrecs; // number of attributes and records
    vector<string> attrnames; // schema
    vector<int> attrinds; // mapping schema to indices
    list<Record*> recs; // list of records
    // methods
    ...
    ...
};
```

Relational algebra (https://en.wikipedia.org/wiki/Relational_algebra) describes 5 fundamental operations, each creating a new relation object from one or two. These operations can be used to build more complex transformations on relational data.

1. Union: All records of both R1 and R2.
Relation * **union**(Relation * R1, Relation * R2)
2. Difference: Records in R1 but not in R2
Relation * **difference** (Relation * R1, Relation * R2)
3. Cartesian Product: All possible pairs of records from R1 and R2
Relation * **cartesianproduct** (Relation * R1, Relation * R2)

4. Projection: New relation with a subset of columns
`Relation * projection (Relation * R1, list<string> projectattrs)`
5. Selection: New relation with a subset of records matching a boolean expression in the attribute values in disjunctive normal form.
`Relation * union (Relation * R1, DNFformula * f)`
 where
`struct DNFformula { list <list <tuple <string, char, Attr *> >`
`> ops; }` is a list of list of tuples with each tuple representing a comparison.
 The top level list stores disjunctions of clauses, each of which represents a list of conjunctions of comparisons. For simplicity assume that there are no negations.
6. Rename: rename an attribute in schema
`Relation * difference (Relation * R1, string s1, string s2)`

Task 1 (50 marks): Implement the above 6 functionalities assuming that the generic comparison operators (`==`, `!=`, `<`, `<=`, `>`, `>=`) have been overloaded for each `Attr` object. Although the prototypes given here are non-member functions, implement the functionalities and member functions, when appropriate.

Task 2 (40 marks): Implement some common Attribute types, e.g.

`integerAttribute`, `floatAttribute`, `stringAttribute`, etc. and overload the operators appropriately. Also, implement the natural join operator using the above implemented primitive operations:

`Relation * naturaljoin (Relation * R1, Relation * R2, list<string> joinattr)`

The list of names in `joinattr` should be attribute names in both `R1` and `R2`, and are the join attributes. See the excerpt below for hints on how to implement (source: https://en.wikipedia.org/wiki/Relational_algebra).

The natural join can be simulated with Codd's primitives as follows. Assume that c_1, \dots, c_m are the attribute names common to R and S , r_1, \dots, r_n are the attribute names unique to R and s_1, \dots, s_k are the attribute names unique to S . Furthermore, assume that the attribute names x_1, \dots, x_m are neither in R nor in S . In a first step the common attribute names in S can be renamed:

$$T = \rho_{x_1/c_1, \dots, x_m/c_m}(S) = \rho_{x_1/c_1}(\rho_{x_2/c_2}(\dots \rho_{x_m/c_m}(S) \dots)) \quad (2)$$

Then we take the Cartesian product and select the tuples that are to be joined:

$$P = \sigma_{c_1=x_1, \dots, c_m=x_m}(R \times T) = \sigma_{c_1=x_1}(\sigma_{c_2=x_2}(\dots \sigma_{c_m=x_m}(R \times T) \dots)) \quad (3)$$

Finally we take a projection to get rid of the renamed attributes:

$$U = \Pi_{r_1, \dots, r_n, c_1, \dots, c_m, s_1, \dots, s_k}(P) \quad (4)$$

Task 3 (10 marks): Provide a simple console based user interface which provides the following functionalities using the above developed “library”:

- Create a new empty table: the user can interactively enter the schema.
- Delete an existing table with all data in it.
- Add a record to a table.
- Print a table in appropriate format.
- Create a table from existing tables using the above developed operations.

Note:

70% marks are for implementing the above functionalities correctly and efficiently. 20% marks are for adhering to the object oriented programming paradigms (e.g. type safety, encapsulation, etc.), and 10% marks are for well-structured code with good documentation.

Submission:

Divide the code into the following:

- A header file (rdb.h) which has all the Library API.
- A c++ file (rdb-basics.cpp) with implementations of all the basic operators.
- A c++ file (rdb-join.cpp) with implementation of the join operation.
- A c++ file (rdb-attr.cpp) with implementations of all supported attribute types.
- A c++ file (rdb-main.cpp) with implementation of the UI.

Submit a zip file containing all the code files for the above problem in Moodle.